October 21, 2026

Manuscript ID of the first submission: Access-2026-44920
Revised title: Does Rubric-Guided Refinement Improve Automatically Optimized Prompts? A Pre-Specified Negative-Result Study
Manuscript type: Negative Result
Corresponding author: Quang-Vinh Dang (British University Vietnam), vinh.dq4@buv.edu.vn

Dear Editor,

Thank you for the opportunity to resubmit. The revised manuscript is a new study. It re-tests Meta-Prompted Instruction Refinement (MPIR), a rubric-guided layer that refines the output of automatic prompt optimization, after addressing each reviewer concern with new experiments and with a reframing of the paper as a negative-result report.

What changed. We repaired the refinement loop, whose rounds were not independent attempts, logged the outcome of every round, and ran 45 base prompts (three optimizers, nine Big-Bench Hard tasks, up to three seeds) under decision rules fixed in a time-stamped protocol before the new experiments. We added a serving-noise measurement with a minimum detectable effect, strict and lenient answer extraction, oracle scoring of all candidates, a gated rubric-free control, format-protected variants, a refiner-capability probe, a judge-validity analysis against a stronger reference judge, and a catalogue of pipeline defects. We no longer call the pilot pre-registered; the protocol is version-controlled, not entered in a registry, and the paper says so.

What we found. With a 1.7B model refining its own prompts the layer changed the prompt in 1 of 80 runs and won in none of 45 cells. A much larger refiner obtained gains under strict scoring in several cells, but most of them reverse under lenient scoring and a rubric-free rewrite reaches gains of the same kind, so the rubric is not shown to be what helps. Two cautionary results seem useful beyond this method: identical prompts do not score identically on a common serving stack, and a gate based on strict answer extraction can accept rewrites that repair the answer tag while harming reasoning.

What we could not do. We could not rerun the GPT-4o and GPT-3.5-turbo configuration of the first evaluation, and we collected no human ratings of the rubric; the paper says so and treats the first evaluation's numbers as exploratory (reproduced, with the validation leak flagged, in an appendix).

For transparency: an earlier version of the first evaluation was submitted to another journal and was not externally reviewed. It has not been published or posted as a preprint. This manuscript is not under review elsewhere, and all authors have approved this submission.

All code, configuration files, the revision protocol with its dated amendments, per-round logs, generations and analysis outputs are available at https://github.com/millennials92/meta_prompted_instruction_refinement.

Thank you for your consideration.

On behalf of all authors,
Quang-Vinh Dang