November 3, 2026

Manuscript ID: Access-2026-54235 (initial submission: Access-2026-44920)
Revised title: Does Rubric-Guided Refinement Improve Prompts from Small-Model Optimizers?
Manuscript type: Negative Result
Corresponding author: Minh Ngoc Dinh (Millenia Education), Minh.dinh@maeducation.com

Dear Editor,

Thank you for the opportunity to resubmit. The reviews of Access-2026-54235 identified one central problem: the primary treatment dropped the answer format in almost all candidates, so the null result could not show that rubric-guided refinement fails. The revised manuscript addresses it with new experiments and a narrower scope.

What changed. (i) A positive control: the validation gate accepted a known-good candidate in 24 of 26 cells in which it helps. (ii) A re-run of the main grid with a format- and example-preserving treatment on 44 cells: candidates were accepted in 15 runs, but under a lenient metric the final prompts won in 1 cell and lost significantly in 3. (iii) A validation of the lenient parser on 400 labelled responses (precision 98.4%, recall 96.8%); the labels and the audit of 80 sampled echo and fallback rounds were produced by a language-model assistant, not by human raters, and the manuscript says so. (iv) A rubric-free control under the same protection: a plain rewrite request gained 3.8 points on average under lenient scoring (11 winning cells of 45), while the rubric-guided variant did not, and strong refiners (Qwen3-235B, GPT-6 Luna) improved the prompts under the protected treatment (7.3 and 3.4 points, nine cells, one run each), but a rubric-free rewrite reached similar gains; so the rubric is not shown to help. (v) Exact acceptance bounds (5.8% one-sided for 1 acceptance in 80 runs; the earlier 7% was the two-sided value), a paired comparison of the rubric-guided and rubric-free variants, a single-session rescoring of the exploratory prompts, an intermediate open-weight refiner (Qwen3-8B, 1 of 9 runs changed the prompt), cost per variant, and a list of all base prompts. (vi) The abstract and claims state the scope (Qwen3-1.7B, nine tasks, weak base prompts), and "pre-specified" is replaced by "protocol fixed in advance, with listed amendments".

What we could not do. We have no second target family and no human ratings, and the strong-refiner comparisons have one run per condition. The paper lists these limits.

This manuscript is not under review elsewhere, and all authors have approved this submission. The author list is unchanged, so no byline-change form is enclosed. All code, per-round logs, samples and labels are at https://github.com/millennials92/meta_prompted_instruction_refinement.

On behalf of all authors,
Minh Ngoc Dinh